



Rolling Updates & Multiple Environments



Rolling Updates with Deployments

- Deployment updates are done through changes to the deployment
- A rolling update is initiated automatically if there's more than 1 replica
- Updates are monitored and stopped if pods fail
- An audit log (history) is kept over all changes
- You can manually roll back to any revision in the history



Update Strategies

- There's 2 update strategies:
 - Recreate
 - RollingUpdate
- Rolling updates can be adjusted by 2 variables:
 - `maxUnavailable` defines the number (or percentage) of Pods that can be unavailable during the update process.
 - `maxSurge` specifies the maximum number (or percentage) of Pods that can be created above the desired number of Pods.



Blue/Green & Canary Deployments

Blue/Green & Canary Deployments can be implemented in two ways:

- With `kubectl rollout pause` and `kubectl rollout resume`
- With a temporary second Deployment (with fewer replicas)



Multiple Environments & Isolation

How many environments do you need?

- How many stages?
- How many teams?

What level of isolation do you need?



Isolation by Namespaces

Default:

- Isolation is just a default to own namespace
- DNS access over namespaces still possible

Additional Isolation:

- Network Policies
- Admission Rules & RBAC
- Isolation by Nodes (against resource competition)



Isolation by Clusters

- Highest level of isolation
- Management Overhead
- Resource/Cost Overhead
- Spinning up new clusters slower than new namespaces
- Switching clusters a bit harder
- Needed for certain (enterprise) compliance & security requirements



Multi-Environment Best Practices

Keep Configuration out of Images and Deployments -> Config Maps & Secrets

Abstract from actual implementation -> Services